

Program Instrumentation with QEMU

RCFC'4

2011/12/04

christophe.guillon@st.com

cedric.vincent@st.com

Outline

Context

Motivations

Program instrumentation

Dynamic binary translation

Execution time and translation time instrumentation

Emulation and instrumentation overhead

Contributions and future work

Context

QEMU, a versatile opensource tool used for:

Platform emulation	(Google SDK,...)
Devices emulation	(VirtualBox,...)
Program emulation	(Scratchbox,...)

Our context:

Emulation of Linux programs (ARM, SH4, ST200, x86)

Our focus:

TCG (Tiny Code Generator), the QEMU compiler

Motivations

Initial motivations:

Program performance analysis

cycles count, profiles, call graphs, ...

Profile driven compiler optimizations

l-cache placement, edge profiling, data dep. estimation, ...

Other usages:

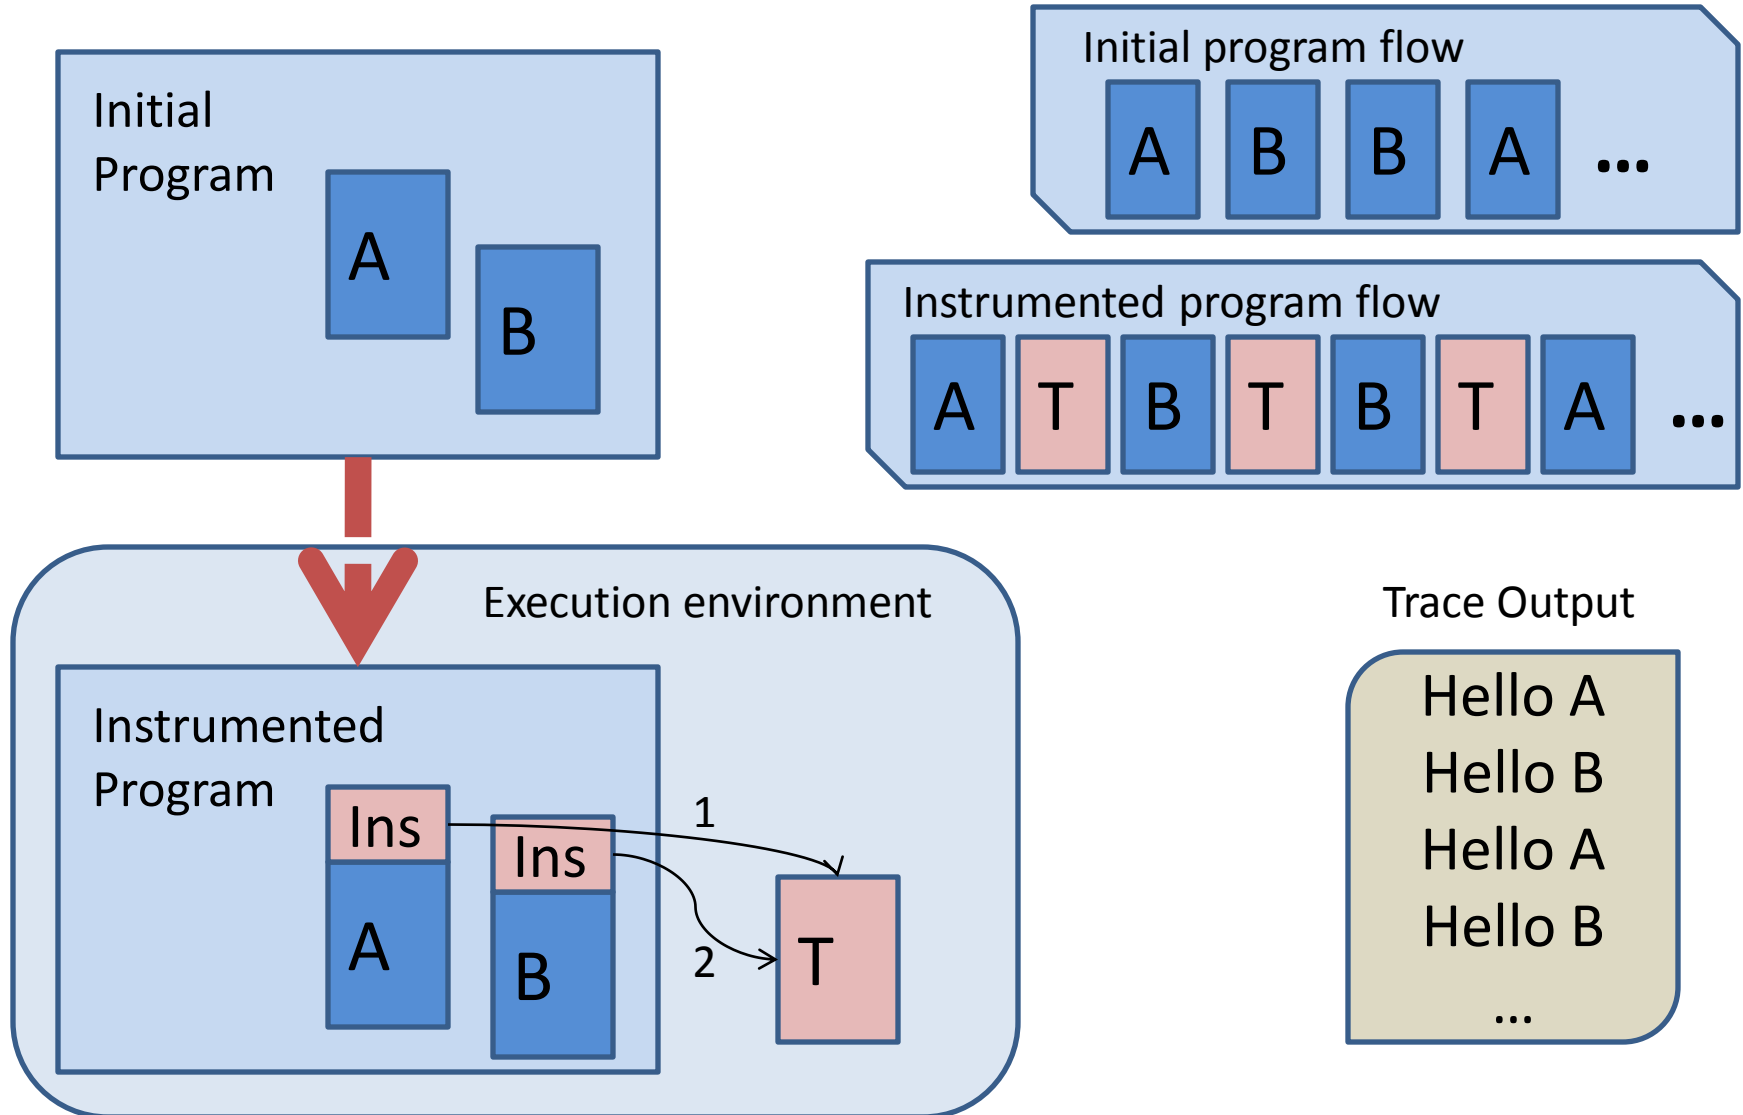
Program debugging

call traces, syscall traces, memory checks, ...

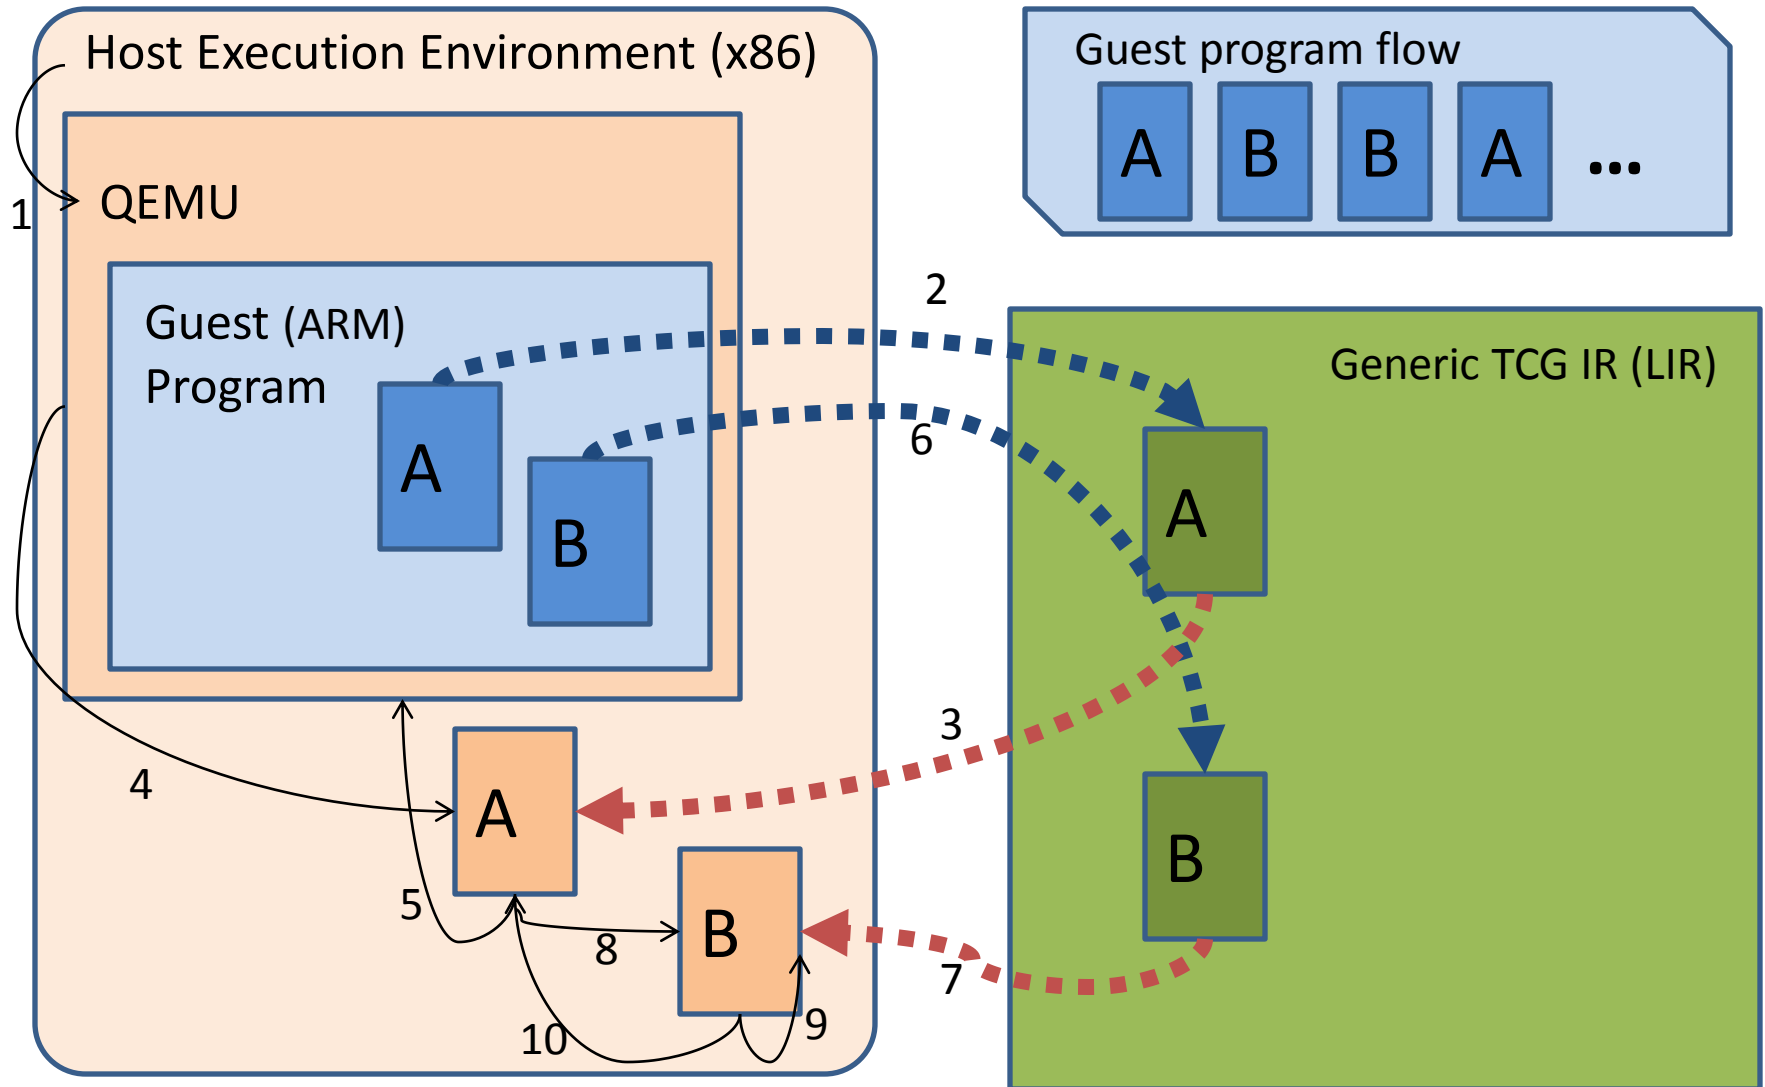
Processor architecture analysis

instructions usage, cache behavior, ...

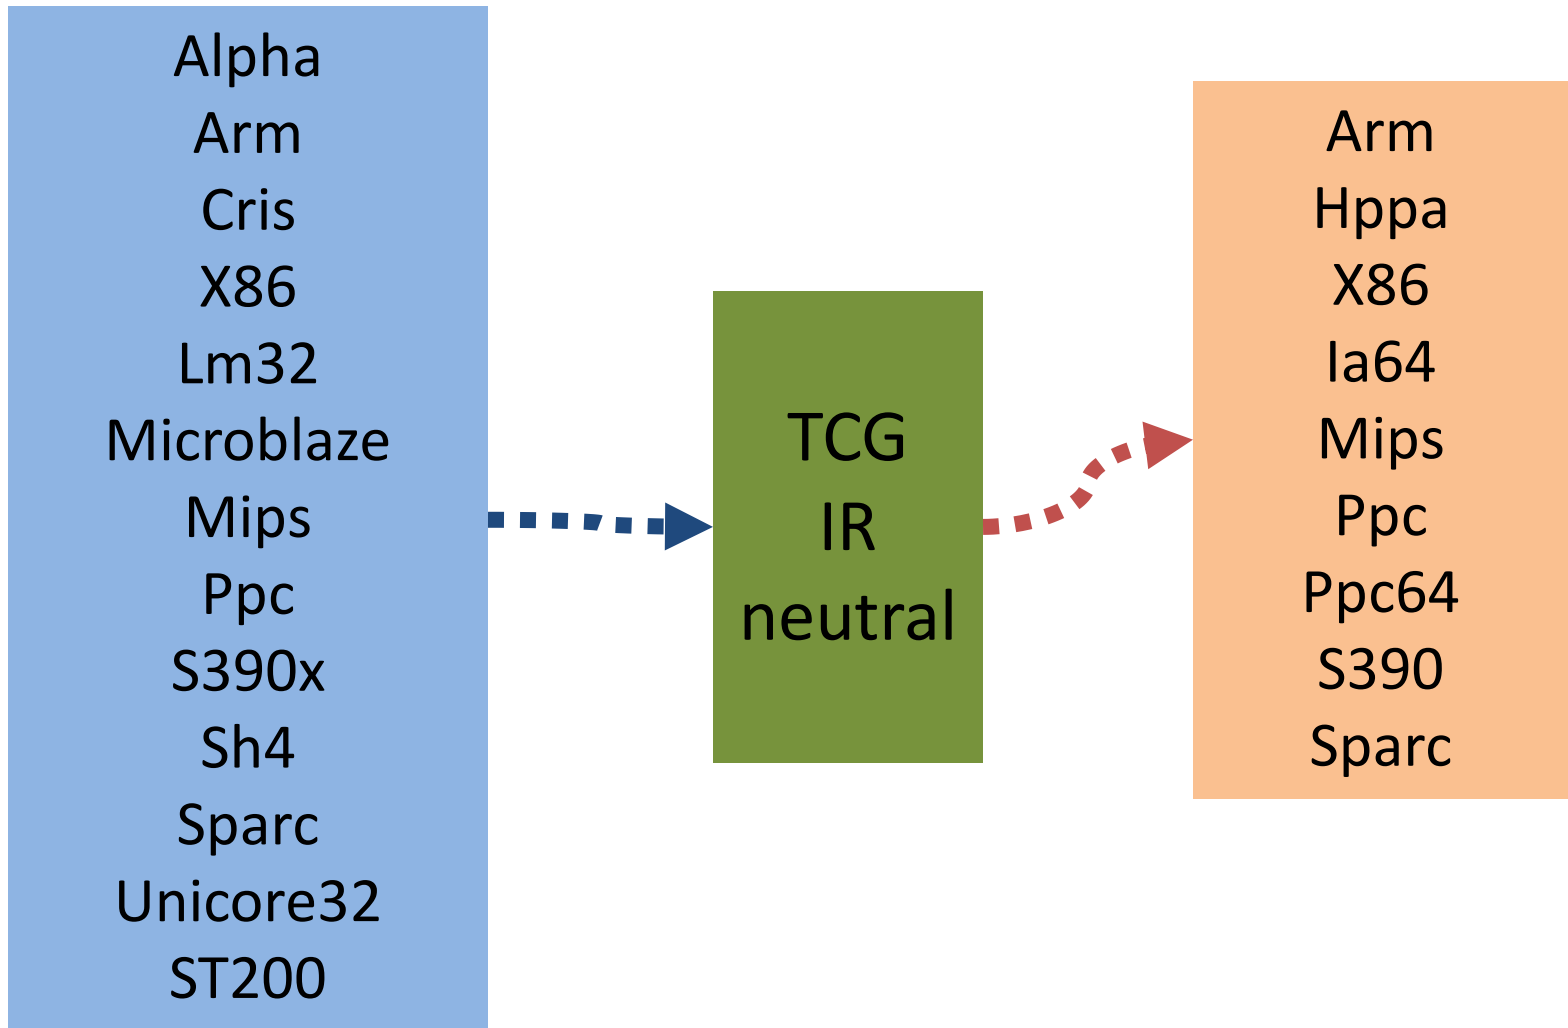
Program Instrumentation




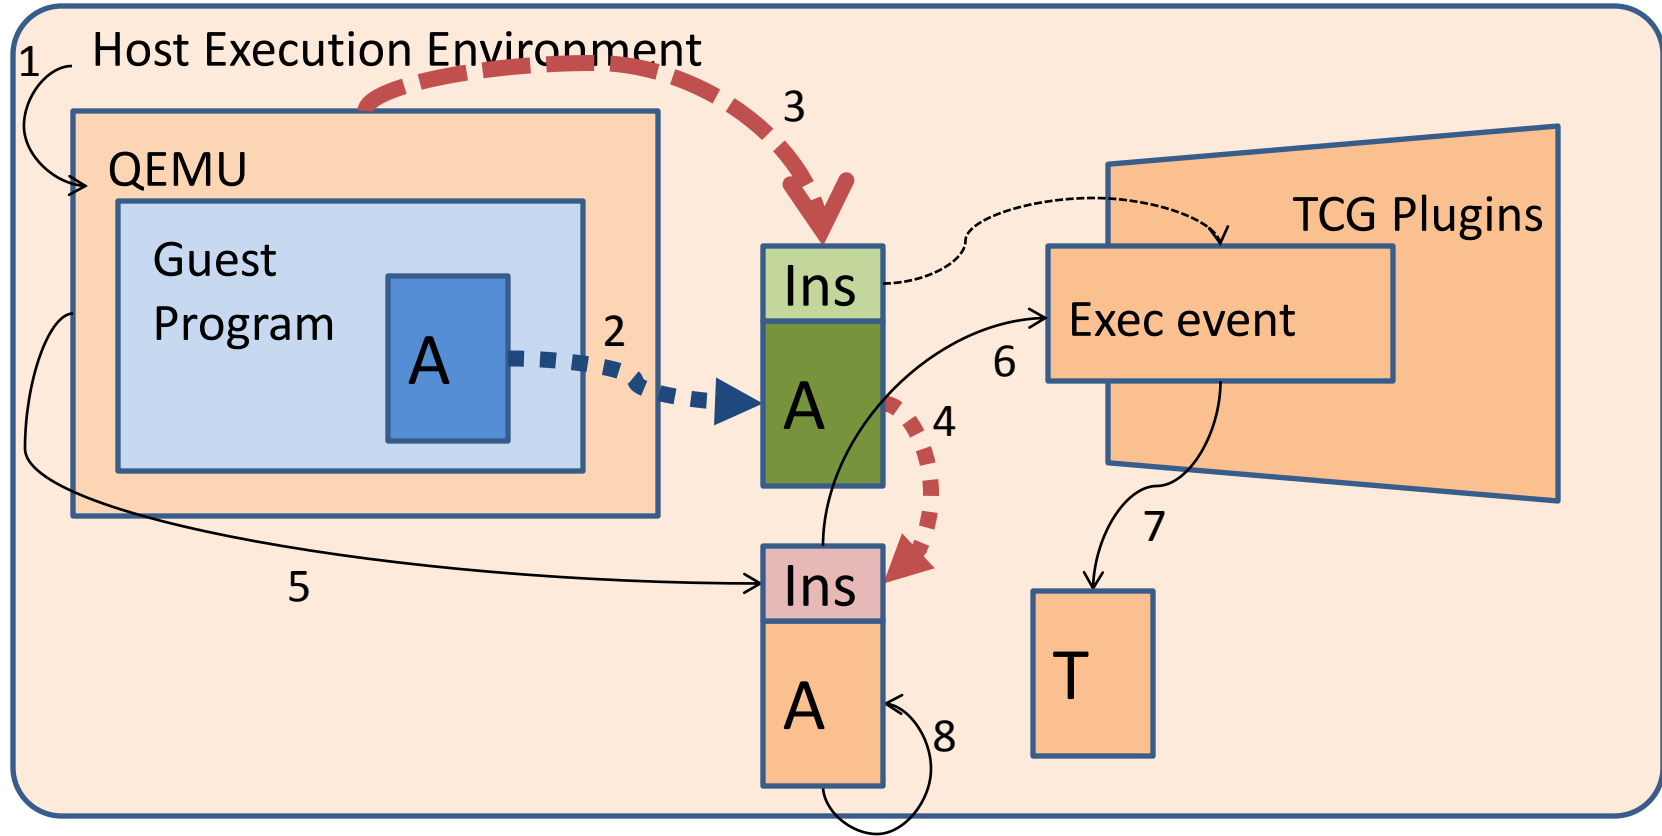
Binary Translation (i.e. QEMU ARM -> x86)



QEMU guests & hosts



Execution Time Instrumentation



Instruction Count Plugin

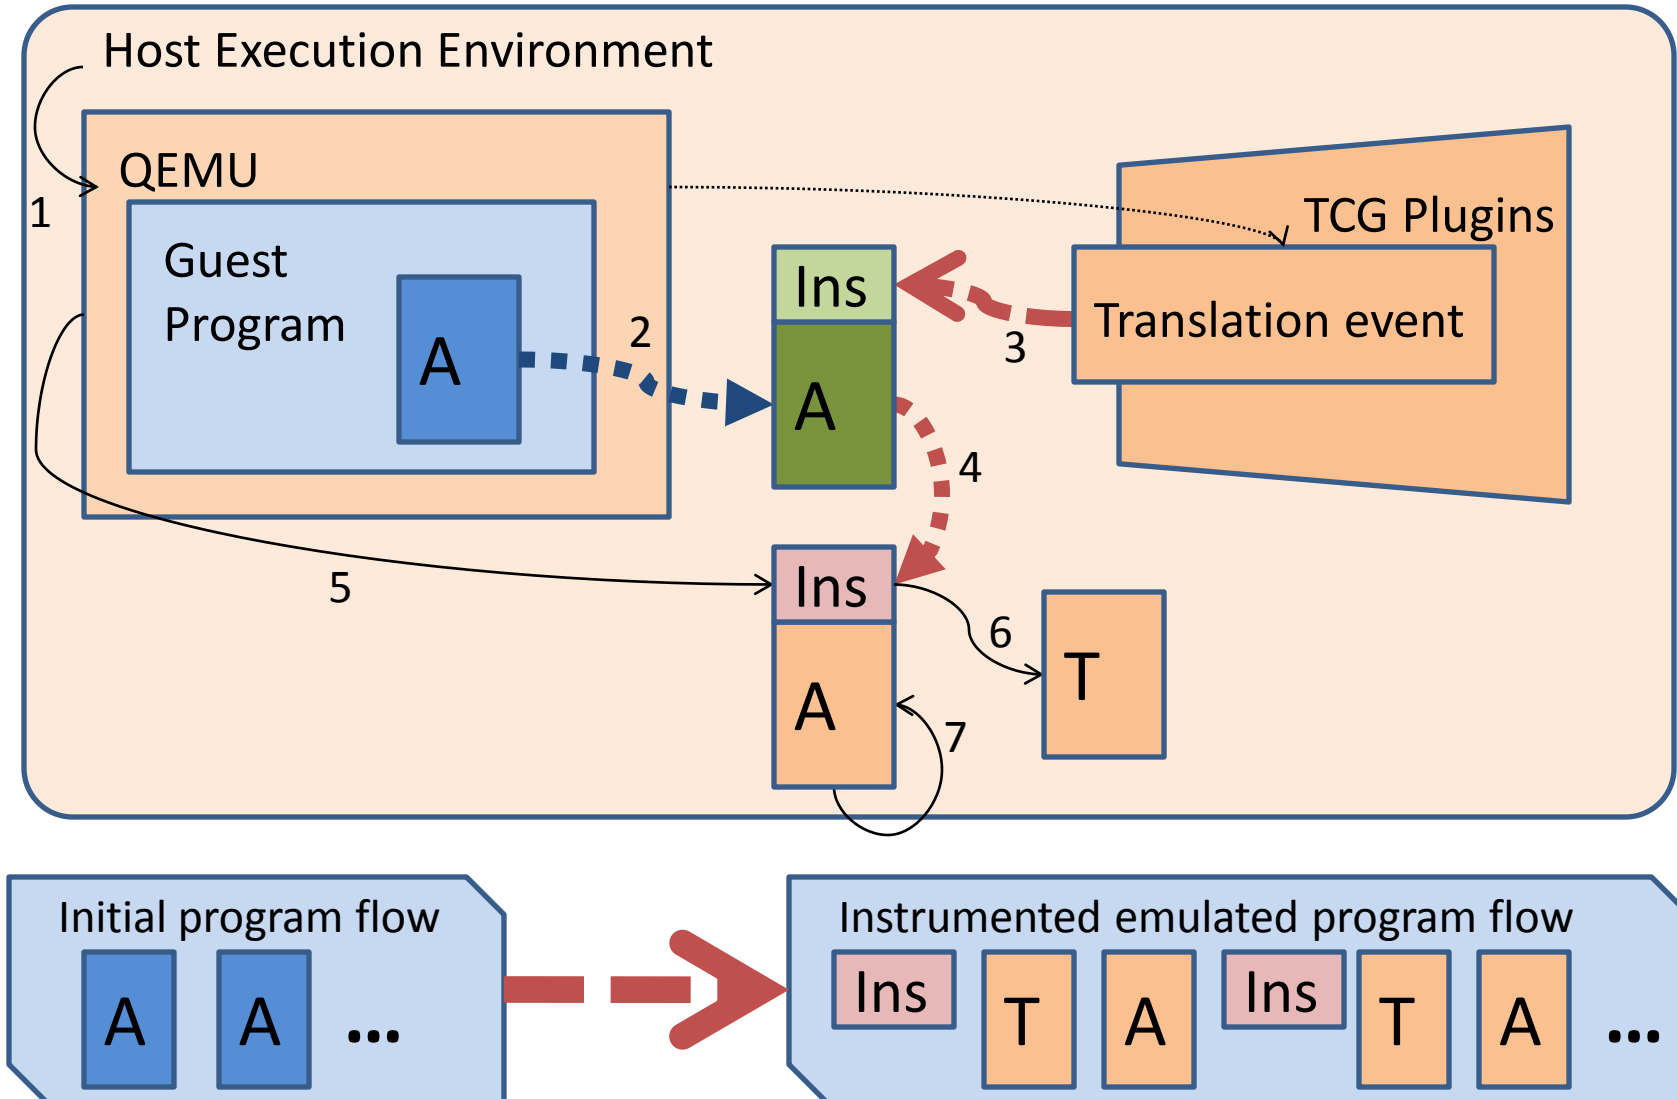
```
/* A simple plugin for counting executed guest instructions.
 * Usage : cc shared -o icount.so icount.c
 * qemu-i386 -tcg-plugin ./icount.so the_program
 */
#include <stdint.h>
#include <stdio.h>
#include « qemu_tpi.h »

uint64_t total_icount;

/* Instruction count updated at each block execution. */
void qemu_tpi_block_execution_event(qemu_tpi_t *tpi) {
    total_icount += TPI_tb_icount(tpi);
}

void qemu_tpi_fini(qemu_tpi_t *tpi) {
    printf("Instructions: %"PRIu64"\n", total_icount);
}
```

Translation Time Instrumentation



Inlined Instruction Count Plugin

```
/* The translation event based code for the icount.c plugin. */

void qemu_tpi_block_translation_event(qemu_tpi_t *tpi) {
    /* Instruction count update inlined at each block translation.
     * Code is generated into the TCG IR buffer directly. */
    TPIv_ptr ptr = TPI_const_ptr(tpi, &total_icount);
    TPIv_i64 total = TPI_temp_new_i64(tpi);
    TPI_gen_ld_i64(tpi, total, ptr, 0);
    TPIv_i64 icount = TPI_const_i64(tpi, TPI_tb_icount(tpi));
    TPI_gen_add_i64(tpi, total, total, icount);
    TPI_temp_free_i64(tpi, icount);
    TPI_gen_st_i64(tpi, total, ptr, 0);
    TPI_temp_free_i64(tpi, total);
    TPI_temp_free_ptr(tpi, ptr);
}
```

TCG Plugins Command Line

```
$ qemu-i386 -tcg-plugin icount.so sha1-i386.exe  
SHA1=15dd99a1991e0b3826fede3deffc1feba42278e6  
Instructions: 107270808
```

```
$ proot -Q qemu-sh4 -tcg-plugin icount.so /root-sh4/ sha1-sh4.exe  
SHA1=15dd99a1991e0b3826fede3deffc1feba42278e6  
Instructions: 166854038
```

```
$ proot -Q qemu-arm -tcg-plugin icount.so /root-arm/ sha1-arm.exe  
SHA1=15dd99a1991e0b3826fede3deffc1feba42278e6  
Instructions: 95419722
```

proot: a sandboxing tool developed at STM


Profile Plugin Example

```
$ qemu-i386 -tcg-plugin profile.so sha1-386.exe
```

```
SHA1=15dd99a1991e0b3826fede3deffc1feba42278e6
```

Instrs	bytes	blocks	symbol
106497664	364229691	1792028	SHA1Transform
571297	1815640	133175	SHA1Update
961	3399	168	SHA1Final
23590	85199	6141	main
18	40	3	__libc_csu_init
16	43	5	.init
2093	12552	2087	.plt
55	146	10	.text
12	28	3	.fini
175102	504418	50840	unknown/libc.so.6

Instrumentation Overhead



Emulation cost is x15 (qemu)

Instrumentation overhead is 30% (offline) for a I-count plugin

Overhead reduced to 3% with translation time interface (inline)

Contributions

Plugin interface for program instrumentation with QEMU

Per code block and per instruction instrumentation

Execution and translation time events

Access to TCG code generator API

Program and libraries symbol table access

Several plugins

I/D-Cache, I-Count, Profile, PC-sample, IO-mem

Companion tool

PRoot sandboxing tool for easy use in Linux user-mode emulation

Future work

Push to QEMU mainstream

Easier access to debug information

Memory checker plugin

Delinquant load detection (compiler feedback)

Gcov like edge-profile (compiler feedback)

Other compilation oriented usage...